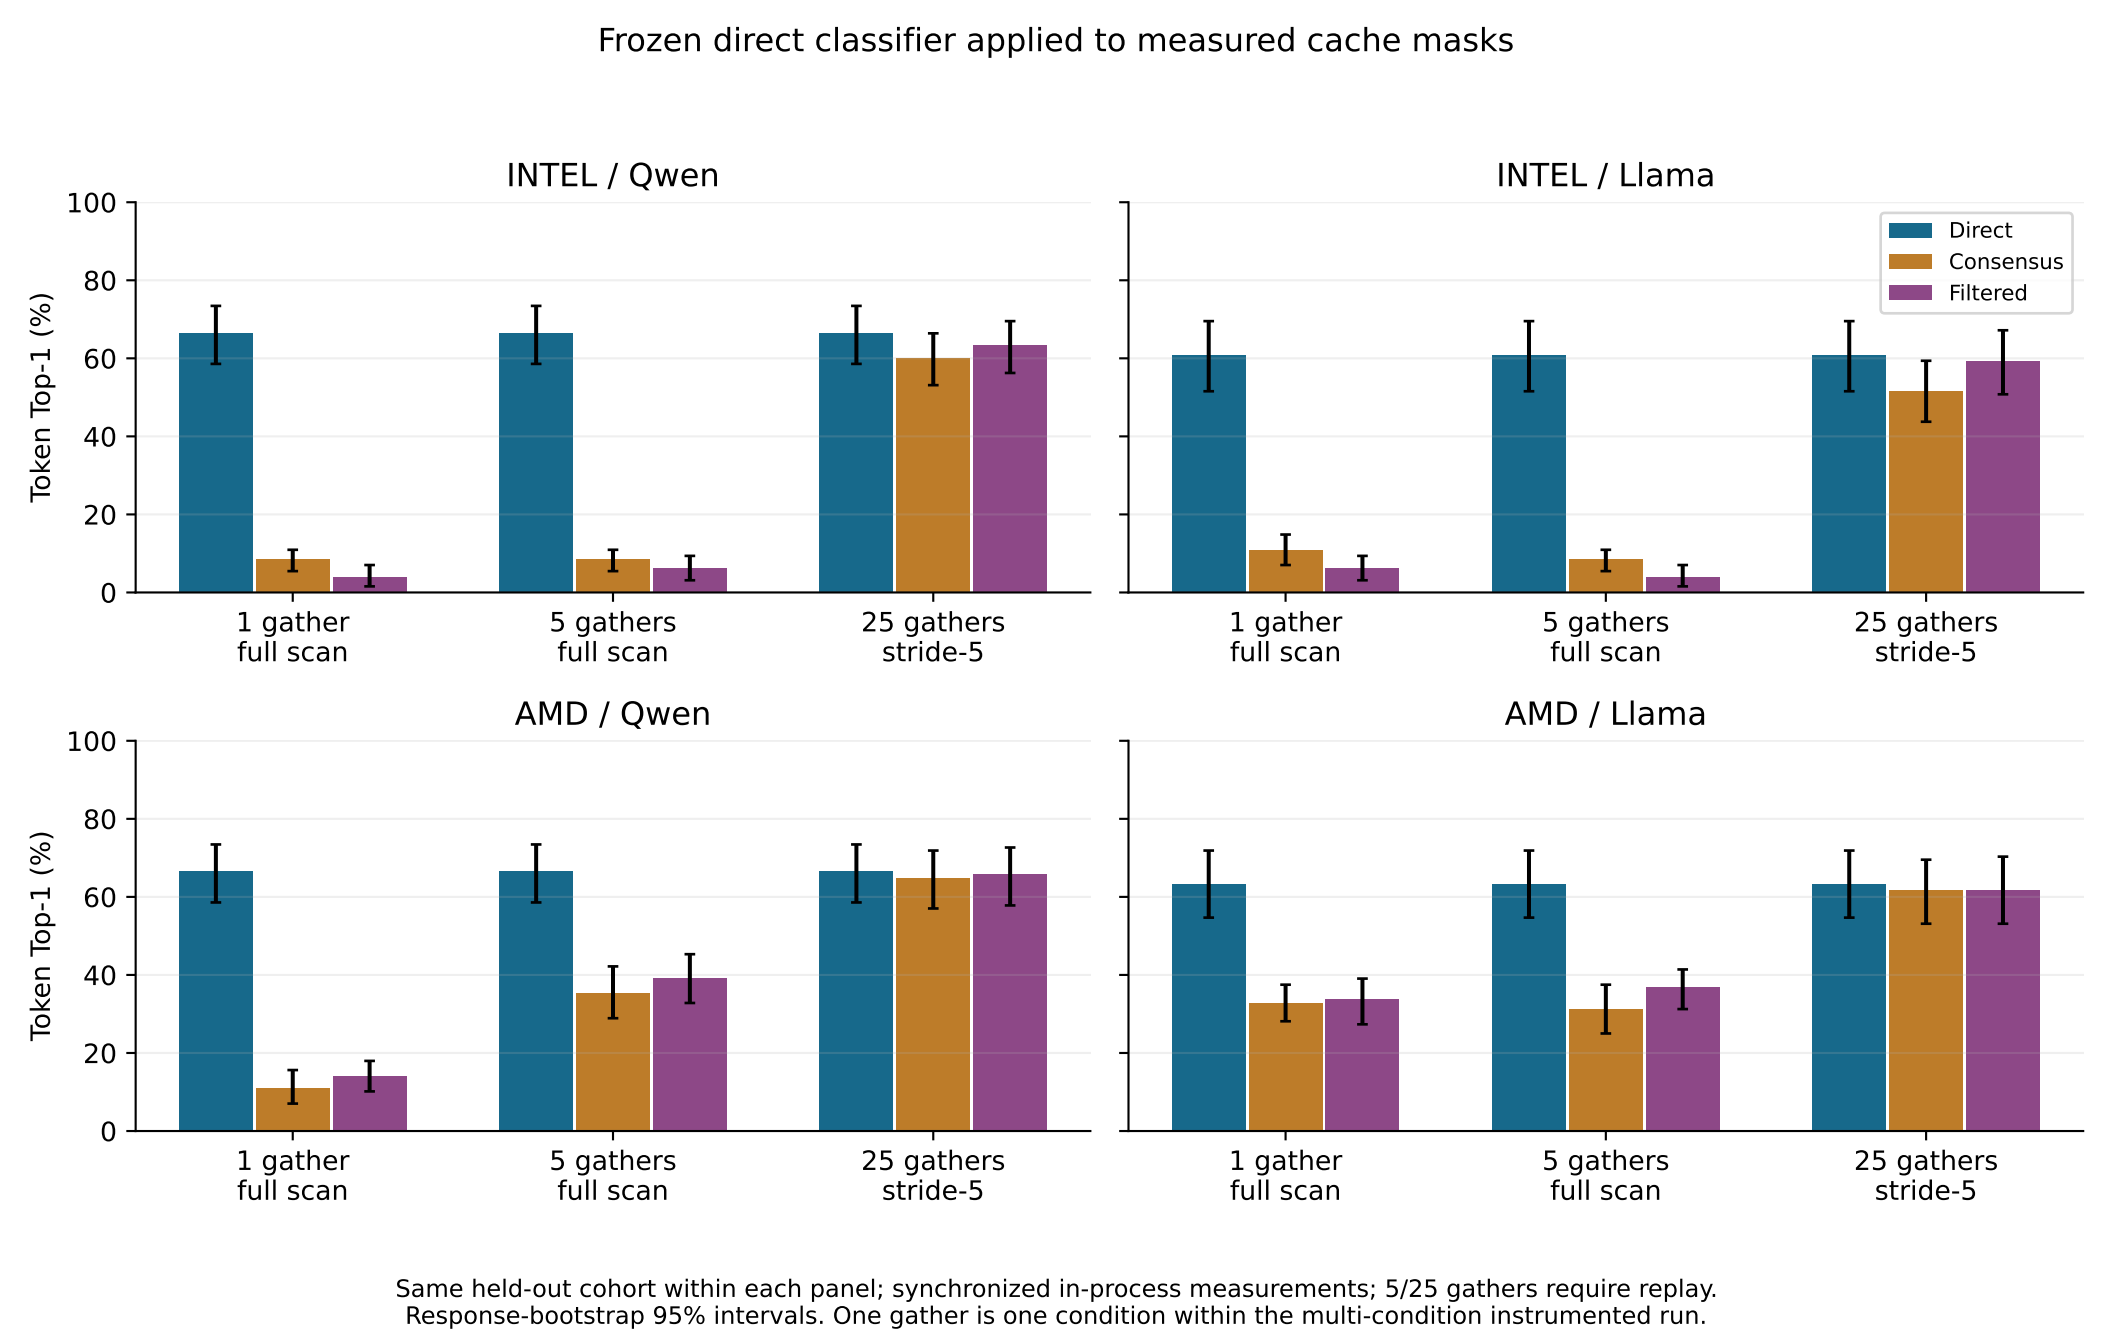
Frozen direct classifier applied to measured cache masks
INTEL / Qwen
INTEL / Llama
AMD / Qwen
AMD / Llama
Token Top-1 (%)
0
20
40
60
80
100
Direct
Consensus
Filtered
1 gather
full scan
5 gathers
full scan
25 gathers
stride-5
Same held-out cohort within each panel; synchronized in-process measurements; 5/25 gathers require replay.
Response-bootstrap 95% intervals. One gather is one condition within the multi-condition instrumented run.